


Weekly Assignment 8

Gaspard Gaches (4645251)

2025-04-22

```
cost = c(.001, .01, .1, 1)
epsilon = c(.001, .01, .1, 1)
set.seed(4645251)
tune.out = tune(svm, medv ~ ., data=Boston[train,], kernel='linear',
               , ranges=list(cost=cost, epsilon=epsilon))

perf = tune.out$performances
perf$cost = factor(perf$cost)
ggplot(perf) + geom_line(aes(epsilon, error, group=cost, color=cost)) + scale_x_log10()
```



```
tune.out$best.parameters
```

```
## cost epsilon
## 4    1    0.001
```

```
tune.out$performances
```

##	cost	epsilon	error	dispersion
## 1	0.001	0.001	43.29272	13.078277
## 2	0.010	0.001	26.95169	10.557590
## 3	0.100	0.001	25.60440	10.829528
## 4	1.000	0.001	25.23923	10.888296
## 5	0.001	0.010	43.30323	13.053327
## 6	0.010	0.010	26.97850	10.603666
## 7	0.100	0.010	25.52956	10.792008
## 8	1.000	0.010	25.37565	11.032447
## 9	0.001	0.100	43.28249	12.821787
## 10	0.010	0.100	26.38916	10.443813
## 11	0.100	0.100	25.50024	10.688779
## 12	1.000	0.100	25.37825	10.714303
## 13	0.001	1.000	56.78102	9.859582
## 14	0.010	1.000	44.25528	7.906997
## 15	0.100	1.000	32.87987	8.671015
## 16	1.000	1.000	30.23543	6.873077

The best parameter tuple (cost, epsilon) from this 10-fold cross-validation training is (1, 0.001). The corresponding purple curve is convex throughout the entire range. The detailed performance results were also printed: we see that ($cost = 1$, $\epsilon = 0.001$) indeed has the lowest error, closely followed by ($cost = 1$, $\epsilon = 0.01$). The ($cost = 0.001$, $\epsilon = 0.001$) pair produced the highest error. Ultimately, a very low margin of tolerance and a high cost yields the lowest cross-validation MSE.

```
best.model = tune.out$best.model
ypred = predict(best.model, Boston[test,])
ytrue = Boston$medv[test]

mse = mean((ytrue - ypred)^2)
mae = mean(abs(ytrue - ypred))
cat('MSE:', round(mse, 4), '\n')
```

```
## MSE: 26.2043
```


```
cat('MAE:', round(mae, 4))
```

```
## MAE: 3.5494
```

Let us now try a grid search with smaller resolution and compare the test MSEs:

```
cost = c(.005, .01, .025, .05, .1, .5, 1)
epsilon = c(.001, .005, .0075, .01, .025, .05, .1, .5, .75, 1)
set.seed(4645251)
tune.out = tune(svm, medv ~ ., data=Boston[train,], kernel='linear',
               , ranges=list(cost=cost, epsilon=epsilon))

perf = tune.out$performances
perf$cost = factor(perf$cost)
ggplot(perf) + geom_line(aes(epsilon, error, group=cost, color=cost)) + scale_x_log10() + scale_y_log10()
```



```
tune.out$best.parameters
```

```
##      cost epsilon
## 54  0.1      0.5
```

After carrying out a more elaborate grid search, the best parameter tuple turns out to be ($cost = 0.1$, $\epsilon = 0.5$). This is a much more balanced pair, with a much higher margin of tolerance and softer penalisation outside of the margin. The cross-validation MSE obtained with this pair is 22.96781, which is lower than the previous best training MSE from the higher resolution grid.

```
best.model = tune.out$best.model
ypred = predict(best.model, Boston[test,])
ytrue = Boston$medv[test]
```

```
mse = mean((ytrue - ypred)^2)
mae = mean(abs(ytrue - ypred))
cat('MSE:', round(mse, 4), '\n')
```

```
## MSE: 25.0509
```

```
cat('MAE:', round(mae, 4))
```

```
## MAE: 3.6301
```


And with this pair we also obtain a lower test MSE - almost 5% lower, which is not a huge gain. Nevertheless, it is interesting to note that the mean absolute error is slightly higher for this pair than for the previous best pair - 2% higher. The previous higher cost parameter likely lead to a little bit of overfitting.

```
cost = c(.001, .01, .1, 1)
epsilon = c(.001, .01, .1, 1)
gamma = c(.001, .01, .1, 1)
set.seed(4645251)
tune.out = tune(svm, medv ~ ., data=Boston[train,], kernel='radial',
               , ranges=list(cost=cost, epsilon=epsilon, gamma=gamma))

perf = tune.out$performances
perf$cost = factor(perf$cost)
res <- tune.out$performances

# Plot: one for each gamma
ggplot(perf, aes(epsilon, error, group=cost, color=cost)) +
  geom_line() +
  facet_wrap(~ gamma, scales = "free") +
  labs(
    title = "MSE across epsilon and cost for each gamma",
    x = "Epsilon",
    y = "Error",
    color = "Cost"
  ) +
  theme_minimal()
```

MSE across epsilon and cost for each gamma



```
tune.out$best.parameters
```

```
##      cost epsilon gamma
## 36      1    0.001   0.1
```

After performing cross-validation with an SVR with a radial basis kernel, we find that the best parameter triplet is ($cost = 1$, $\epsilon = 0.001$, $\gamma = 0.1$). The first two values are familiar, since they are the exact same obtained with a linear kernel and the same resolution grid in our first experiment.

```
best.model = tune.out$best.model
ypred = predict(best.model, Boston[test,])
ytrue = Boston$medv[test]
```

```
mse = mean((ytrue - ypred)^2)
mae = mean(abs(ytrue - ypred))
cat('MSE:', round(mse, 4), '\n')
```

```
## MSE: 17.1149
```

```
cat('MAE:', round(mae, 4))
```

```
## MAE: 2.442
```

Both the MSE and MAE are significantly lower than what our best parameter pair yielded with a linear kernel.

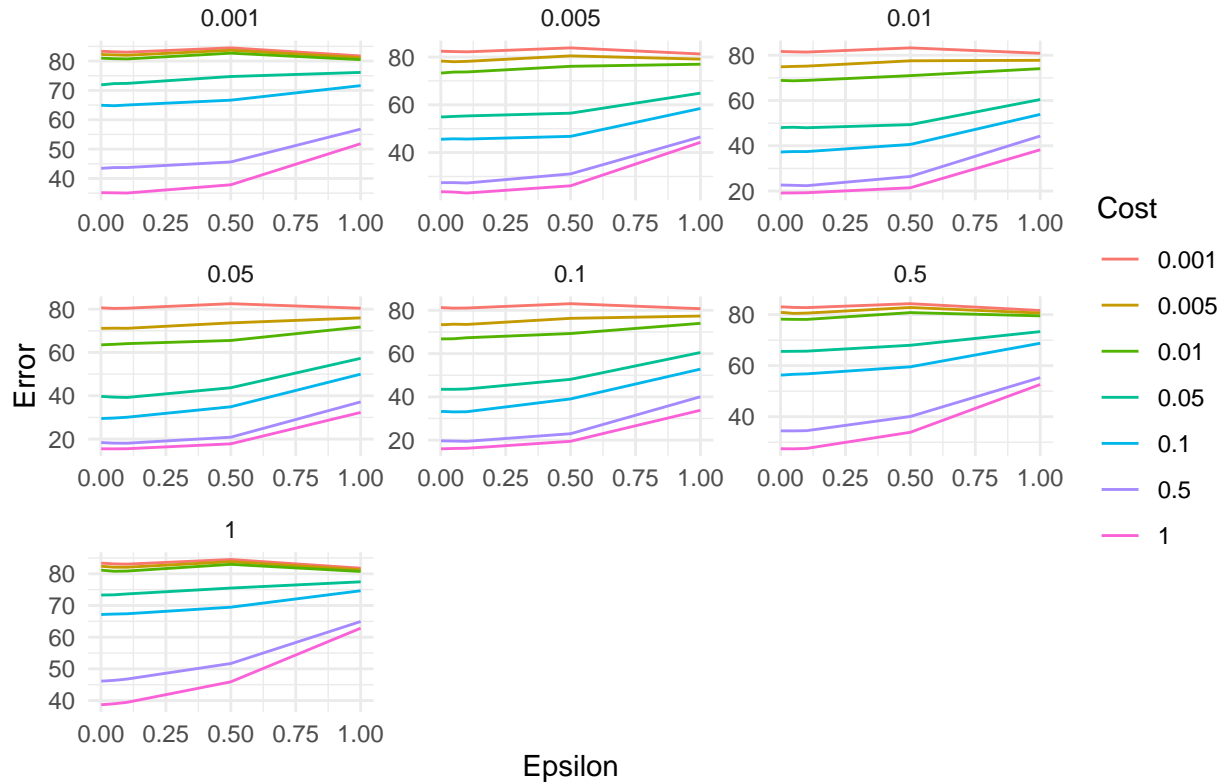
Let us now repeat this experiment with a higher resolution grid:

```
cost = c(.001, .005, .01, .05, .1, .5, 1)
epsilon = c(.001, .005, .01, .05, .1, .5, 1)
gamma = c(.001, .005, .01, .05, .1, .5, 1)
set.seed(4645251)
tune.out = tune(svm, medv ~ ., data=Boston[train,], kernel='radial',
               , ranges=list(cost=cost, epsilon=epsilon, gamma=gamma))

perf = tune.out$performances
perf$cost = factor(perf$cost)
res <- tune.out$performances

# Plot: one for each gamma
ggplot(perf, aes(epsilon, error, group=cost, color=cost)) +
  geom_line() +
  facet_wrap(~ gamma, scales = "free") +
  labs(
    title = "MSE across epsilon and cost for each gamma",
    x = "Epsilon",
    y = "Error",
    color = "Cost"
  ) +
  theme_minimal()
```

MSE across epsilon and cost for each gamma



```
tune.out$best.parameters
```

```
##      cost epsilon gamma
## 175     1    0.05  0.05
```

```
best.model = tune.out$best.model
ypred = predict(best.model, Boston[test,])
ytrue = Boston$medv[test]
```

```
mse = mean((ytrue - ypred)^2)
mae = mean(abs(ytrue - ypred))
cat('MSE:', round(mse, 4), '\n')
```

```
## MSE: 16.8298
```

```
cat('MAE:', round(mae, 4))
```

```
## MAE: 2.4546
```

It turns out that the ($cost = 1$, $\epsilon = 0.05$, $\gamma = 0.05$) triplet gives an even lower MSE - though the gain is really minimal, only 1.5% -, but offers no improvement in terms of MAE. The fact that the best gamma is relatively high means that each point influences a rather small region, so bias is lower here. Our training set is relatively big in statistical terms, so it is no surprise that lowering the bias yields better predictions.

```
lm.model <- lm(medv ~ ., data = Boston[train, ])  
lm.pred <- predict(lm.model, newdata = Boston[test, ])  
  
mse = mean((ytrue - lm.pred)^2)  
mae = mean(abs(ytrue - lm.pred))  
cat('MSE:', round(mse, 4), '\n')
```

```
## MSE: 26.151
```

```
cat('MAE:', round(mae, 4))
```

```
## MAE: 3.7217
```

The error reduction compared to a simple multivariate regression is indeed significant both in terms of MSE and MAE.